

AI Internship Requirements

The AI internship is designed for candidates who can build practical machine learning features in a full-stack product environment.

Required programming skills include Python, TypeScript, basic SQL, and the ability to read API documentation independently.

Required Python libraries include NumPy, Pandas, Scikit-learn, FastAPI, and basic testing tools such as pytest.

Candidates should understand supervised learning, embeddings, vector search, retrieval augmented generation, model evaluation, and data preprocessing.

Tooling and Project Expectations

The intern is expected to build a small AI assistant that uses a local model, document retrieval, chat memory, and optional web search.

The preferred stack is Next.js for the frontend, FastAPI for the backend, Ollama for local LLM inference, ChromaDB for vector search, and SQLite for chat persistence.

The project should include streaming responses, citation display, deterministic routing, and clear error handling when a source cannot answer the question.

Conclusion

A successful AI intern should demonstrate practical engineering judgement, not only model prompting.

The strongest submission combines full-stack implementation, hybrid retrieval, reliable source citation, deterministic tool routing, and evaluation metrics.

The project should prove that the candidate can build a controllable AI workflow that answers only from trusted evidence unless web search is explicitly enabled.